

Sean Xuefeng Du

Department of Computer Sciences
University of Wisconsin–Madison
1210 W Dayton St, Madison, WI 53706
[Homepage](#), [Scholar](#), +1-608-720-4664
xfdu@cs.wisc.edu, xuefengdu1@gmail.com

Research Interests

Reliable machine learning, foundation model reliability, and related applications, specifically I work on:

- Out-of-distribution (OOD) research: designing adaptive and interpretable learning algorithms that help ML models detect and generalize on OOD samples, such as semantic/covariate shifts, adversarial and noisy samples.
- Reliability of foundation models: understanding the blindspots of LLMs and VLMs for improved safeguarding efforts, such as model hallucination and harmful prompt detection.
- Applications: Graph embedding, object detection and segmentation.
- Interdisciplinary research: Biometrics, AI-aided cryo-microscopy/protein structural analysis.

Education

Ph.D. candidate in Computer Sciences Jan. 2021–Present

[University of Wisconsin–Madison](#), Madison, WI

- Ph.D. research in reliable machine learning.
- Advisor: Prof. [Sharon Yixuan Li](#).

B. Eng. in Electrical Engineering

Sept. 2016 – Jun. 2020

[Xi'an Jiaotong University](#), Xi'an, China

- Overall GPA: 91.60/100(3.83/4.0), Rank: 1st/170.

Awards

- [Rising Stars in Data Science](#), September, 2024.
- [Jane Street Graduate Research fellowship](#) (\$50,000 grant), April, 2023.
- NeurIPS Scholar Award, November, 2022.
- CS departmental research fellowship, UW-Madison, September, 2021.
- National Scholarship (2x), Ministry of Education in China, 2017–2018.

Research Experience

Research Intern June 2024 – Sept. 2024

Microsoft Research, Redmond, WA, USA

- Hosts: Dr. [Robert Sim](#), [Jay Stokes](#) and [Reshmi Ghosh](#).
- Research on *Malicious prompt detection for vision language models*

Student Researcher

June 2022 – Sept. 2022

Google Research, Sunnyvale, CA, USA

- Hosts: Dr. [Zizhao Zhang](#), [Ting Chen](#) and [Han Zhang](#).
- Research on *Open-vocabulary object detection with language models*

Research Intern

Mar. 2021 – June 2021

Tencent AI Lab, Shenzhen, China

- Supervised by Dr. [Yu Rong](#) and [Junzhou Huang](#).

- Research on *Robust graph neural networks against noisy labels*.

Research Assistant Oct. 2020 – Jan. 2021
Department of Computer Science, Hong Kong Baptist University, Hong Kong

- Supervised by Dr. [Bo Han](#).
- Research on *Effective network architecture for adversarial robustness*.

Research Engineer Intern Aug. 2020 -Sept. 2020
AI Lab, Bytedance Inc., Beijing, China

- Supervised by Dr. [Changhu Wang](#).
- Research on *Fine-grained image classification*.

Research Intern Dec. 2019 -July 2020
AI Theory Group, Noah’s Ark Lab, Shenzhen, China

- Supervised by Dr. [Hang Xu](#) and [Chenhan Jiang](#).
- Research on *Hybrid supervised panoptic segmentation*.

Student Intern July 2018 – Feb. 2020
Department of Computational Biology, CMU, Pittsburgh, PA, USA

- Supervised by Dr. [Min Xu](#) and Dr. [Haohan Wang](#).
- Research on *Deep learning for cellular electron cryo-tomography analysis*.

Research Assistant Nov. 2018 – Jun. 2020
Intelligent Networks and Network Security Lab, XJTU, Xi’an, China

- Supervised by Dr. [Pinghui Wang](#).
- Research on *Few-shot node classification with meta learning*.

Research Assistant Jun. 2017 - Apr. 2019
Institute of Automatic Control, XJTU, Xi’an, China

- Supervised by Dr. [Dexing Zhong](#).
- Research on *Machine learning for hand-based biometrics*.

Publications See full list in [my google scholar](#) page, I publish under the name “Xuefeng Du”.
* indicates Equal Contribution

43. HaloScope: Harnessing Unlabeled LLM Generations for Hallucination Detection
Xuefeng Du, Chaowei Xiao, Yixuan Li
Advances in Neural Information Processing Systems (NeurIPS) 2024.
42. [When and How does In-distribution Label Help Out-of-distribution Detection?](#)
Xuefeng Du, Yiyu Sun, Yixuan Li
International Conference on Machine Learning (ICML) 2024.
41. [How does Unlabeled Data Provably Help Out-of-distribution Detection?](#)
Xuefeng Du*, Zhen Fang*, Ilias Diakonikolas, Yixuan Li
International Conference on Learning Representations (ICLR) 2024.
40. [Feed Two Birds with One Scone: Exploiting Wild Data for Both Out-of-Distribution Generalization and Detection](#)
Haoyue Bai, Gregory Canal, **Xuefeng Du**, Jeongyeol Kwon, Robert D Nowak, Yixuan Li
International Conference on Machine Learning (ICML) 2023.

39. [Dream the Impossible: Outlier Imagination with Diffusion Models](#)
Xuefeng Du, Yiyu Sun, Jerry Zhu, Yixuan Li
Advances in Neural Information Processing Systems (NeurIPS) 2023.
38. [OpenOOD v1.5: Enhanced Benchmark for Out-of-Distribution Detection](#)
Jingyang Zhang, Jingkang Yang, Pengyun Wang, Haoqi Wang, Yueqian Lin, Haoran Zhang, Yiyu Sun, **Xuefeng Du**, Kaiyang Zhou, Wayne Zhang, Yixuan Li, Ziwei Liu, Yiran Chen, Hai Li
Advances in Neural Information Processing Systems 2023 DistShift (NeurIPS) Workshop.
37. [Non-parametric Outlier Synthesis](#)
Leitian Tao, **Xuefeng Du**, Jerry Zhu, Yixuan Li
International Conference on Learning Representations (ICLR) 2023.
36. [Noise-robust Graph Learning by Estimating and Leveraging Pairwise Interactions](#)
Xuefeng Du, Tian Bian, Yu Rong, Bo Han, Tongliang Liu, Tingyang Xu, Wenbing Huang, Yixuan Li, Junzhou Huang
Transactions on Machine Learning Research (TMLR) 2023
35. [SIREN: Shaping Representations for Detecting Out-of-distribution Objects](#)
Xuefeng Du, Gabriel Gozum, Yifei Ming, Yixuan Li
Advances in Neural Information Processing Systems (NeurIPS) 2022.
34. [OpenOOD: Benchmarking Generalized Out-of-Distribution Detection](#)
Jingkang Yang, Pengyun Wang, Dejian Zou, Zitang Zhou, Kunyuan Ding, Wenxuan Peng, Haoqi Wang, Guangyao Chen, Bo Li, Yiyu Sun, **Xuefeng Du**, Kaiyang Zhou, Wayne Zhang, Dan Hendrycks, Yixuan Li, Ziwei Liu
Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track, 2022.
33. [Unknown-Aware Object Detection: Learning What You Don't Know from Videos in the Wild](#)
Xuefeng Du, Xin Wang, Gabriel Gozum, Yixuan Li
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) 2022, **oral**.
32. [Performance-Aware Mutual Knowledge Distillation for Improving Neural Architecture Search](#)
Pengtao Xie, **Xuefeng Du**
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) 2022.
31. [VOS: Learning What You Don't Know by Virtual Outliers Synthesis](#)
Xuefeng Du, Eric Wang, Mu Cai, Yixuan Li
International Conference on Learning Representations (ICLR) 2022.
30. [Learning Diverse-Structured Networks for Adversarial Robustness](#)
Xuefeng Du, Jingfeng Zhang, Bo Han, Tongliang Liu, Yu Rong, Gang Niu, Junzhou Huang, Masashi Sugiyama
International Conference on Machine Learning (ICML) 2021.
29. [How to save your annotation cost for Panoptic Segmentation?](#)
Xuefeng Du, Chenhan Jiang, Hang Xu, Gengwei Zhang, Zhenguo Li
The AAAI Conference on Artificial Intelligence (AAAI) 2021.
28. [Active learning to classify macromolecular structures in situ for less supervision in cryoelectron tomography](#)
Xuefeng Du, Haohan Wang, Zhenxi Zhu, Xiangrui Zeng, Yi-Wei Chang, Jing

Zhang, Eric Xing, Min Xu
Bioinformatics, 2021

27. [Node Classification on Graphs with Few-Shot Novel Labels via Meta Transformed Network Embedding](#)
Lin Lan, Pinghui Wang, **Xuefeng Du**, Kaikai Song, Jing Tao, Xiaohong Guan
Advances in Neural Information Processing Systems (NeurIPS) 2020.
26. [A deep biometric hash learning framework for three advanced hand-based biometrics.](#)
Huikai Shao, Dexing Zhong, **Xuefeng Du**
IET Biometrics.
25. [Few-shot learning for palmprint recognition via meta-siamese network](#)
Huikai Shao, Dexing Zhong, **Xuefeng Du**, Shaoyi Du, Raymond NJ Veldhuis
IEEE transactions on instrumentation and measurement.
24. [Deep distillation hashing for unconstrained palmprint recognition](#)
Huikai Shao, Dexing Zhong, **Xuefeng Du**
IEEE transactions on instrumentation and measurement.
23. [Cross-domain palmprint recognition via regularized adversarial domain adaptive hashing](#)
Xuefeng Du, Dexing Zhong, Huikai Shao
IEEE Transactions on Circuits and Systems for Video Technology.
22. [Effective deep ensemble hashing for open-set palmprint recognition](#)
Huikai Shao, Dexing Zhong, **Xuefeng Du**
Journal of Electronic Imaging.
21. [Cross-domain palmprint recognition based on transfer convolutional autoencoder](#)
Huikai Shao, Dexing Zhong, **Xuefeng Du**
IEEE International Conference on Image Processing (ICIP) 2019.
20. [Continual palmprint recognition without forgetting](#)
Xuefeng Du, Dexing Zhong, Huikai Shao
IEEE International Conference on Image Processing (ICIP) 2019.
19. [Building an active palmprint recognition system](#)
Xuefeng Du, Dexing Zhong, Huikai Shao
IEEE International Conference on Image Processing (ICIP) 2019.
18. [Low-shot palmprint recognition based on meta-siamese network](#)
Xuefeng Du, Dexing Zhong, Pengna Li
IEEE International Conference on Multimedia and Expo (ICME) 2019.
17. [A hand-based multi-biometrics via deep hashing network and biometric graph matching](#)
Dexing Zhong, Huikai Shao, **Xuefeng Du**
IEEE Transactions on Information Forensics and Security.
16. [Decade progress of palmprint recognition: A brief survey](#)
Dexing Zhong, **Xuefeng Du**, Kuncai Zhong
Neurocomputing.
15. [Classification in Cryo-Electron Tomograms](#)
Ilja Gubins, Gijs van der Schot, Remco C Veltkamp, Friedrich Förster, **Xuefeng Du**, Xiangrui Zeng, Zhenxi Zhu, Lufan Chang, Min Xu, Emmanuel Moebel, Antonio Martinez-Sanchez, Charles Kervrann, Tuan M Lai, Xusi Han, Genki Terashi, Daisuke Kihara, Benjamin A Himes, Xiaohua Wan, Jingrong Zhang,

Shan Gao, Yu Hao, Zhilong Lv, Xiaohua Wan, Zhidong Yang, Zijun Ding, Xuefeng Cui, Fa Zhang
Eurographics 2019.

14. [Open-set Recognition of Unseen Macromolecules in Cellular Electron Cryo-Tomograms by Soft Large Margin Centralized Cosine Loss](#)
Xuefeng Du, Xiangrui Zeng, Bo Zhou, Alex Singh and Min Xu
British Machine Vision Conference (BMVC) 2019, **Spotlight**.
13. [Semi-supervised Macromolecule Structural Classification in Cellular Electron Cryo-Tomograms using 3D Autoencoding Classifier](#)
Siyuan Liu, **Xuefeng Du**, Rong Xi, Fuya Xu, Xiangrui Zeng, Bo Zhou and Min Xu
British Machine Vision Conference (BMVC) 2019.
12. [Efficient Deep Palmprint Recognition via Distilled Hashing Coding](#) Huikai Shao, Dexing Zhong and **Xuefeng Du**
IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops 2019.
11. [Palm vein recognition with deep hashing network](#)
Dexing Zhong, Shuming Liu, Wenting Wang, **Xuefeng Du**
Pattern Recognition and Computer Vision 2019.
10. [Palmprint recognition using siamese network](#)
Dexing Zhong, Yuan Yang, **Xuefeng Du**
13th Conference of Biometric Recognition, CCBR 2018.

Preprints

9. [How Reliable Is Human Feedback For Aligning Large Language Models?](#)
Min-Hsuan Yeh, Leitian Tao, Jeffrey Wang, **Xuefeng Du**, Yixuan Li.
arXiv preprint arXiv:2410.01957
8. [VLMGuard: Defending VLMs against Malicious Prompts via Unlabeled Data](#)
Xuefeng Du, Reshmi Ghosh, Robert Sim, Ahmed Salem, Vitor Carvalho, Emily Lawton, Yixuan Li, Jack W. Stokes.
arXiv preprint arXiv:2410.00296.
7. [Out-of-Distribution Learning with Human Feedback](#)
Haoyue Bai, **Xuefeng Du**, Katie Rainey, Shibin Parameswaran, Yixuan Li
arXiv preprint arXiv:2408.07772.
6. [The Ghanaian NLP Landscape: A First Look](#)
Sheriff Issaka, Zhaoyi Zhang, Mihir Heda, Keyi Wang, Yinka Ajibola, Ryan DeMar, **Xuefeng Du**
arXiv preprint arXiv:2405.06818.
5. [Exploring Transition States of Protein Conformational Changes via Out-of-Distribution Detection in the Hyperspherical Latent Space](#)
Bojun Liu, Jordan G Boysen, Ilona Christy Unarta, **Xuefeng Du**, Yixuan Li, Xuhui Huang
Chemrxiv, 2024
4. [Learning by Passing Tests, with Application to Neural Architecture Search](#)
Xuefeng Du, Pengtao Xie
arXiv preprint arXiv:2011.15102.
3. [Skilllearn: Machine Learning Inspired by Humans' Learning Skills](#)
Pengtao Xie, **Xuefeng Du**, Hao Ban
arXiv preprint arXiv:2012.04863.

2. [Small-Group Learning, with Application to Neural Architecture Search](#)
Xuefeng Du, Pengtao Xie
arXiv preprint arXiv:2012.12502.
1. [Towards efficient unconstrained palmprint recognition via deep distillation hashing](#)
Huikai Shao, Dexing Zhong, **Xuefeng Du**
arXiv preprint arXiv:2004.03303.

Patents

- VisionGuard: Detecting Multimodal AI Systems from Indirect Prompt Injections (pending)
Xuefeng Du, Reshmi Ghosh, Vitor Carvalho, Robert Sim, Emily Lawton, Jay Stokes, Lukas Wutschitz, Reshmi Ghosh, Ahmed Salem
- [A kind of cross-platform palm grain identification method](#)
Dexing Zhong, Huikai Shao, Xuefeng Du, Runzhao Yao
- [A kind of vehicle early warning vehicle intelligent system based on multisensor](#)
Dexing Zhong, Huikai Shao, Xuefeng Du

Talks

- Talk at Microsoft, 08/2024
- Talk at Jane Street, 04/2023
- Talk at [AI talks](#), 02/2023
- Talk at [MLOPT Seminar](#), UW-Madison, 05/2022
- Talk at Google, 01/2022
- Talk at Adobe, 12/2021
- Talk at Microsoft, 12/2021

Teaching

Teaching Assistant for

- [CS762](#) (graduate course), UW-Madison: Advanced Deep Learning, Fall 2022.
–Designing lectures on topics in trustworthy machine learning, grading students’ written papers, holding office hours, leading discussions and answering questions.
- [CS540](#) (undergraduate course), UW-Madison: Intro to AI, Spring 2021.
–Designing lectures on topics in unsupervised learning, deep learning, and reinforcement learning. Preparing all the homework related to the topic of deep learning, and design final exams, guiding students on final projects, holding office hours, and delivering discussion sections.

Research Mentoring

9. [Seongheon Park](#), PhD student at UW-Madison, present.
8. [Sheriff Issaka](#), undergraduate student at UW-Madison, 2024, now data scientist at Zafar & ICU Data Science Labs.
7. [Leitian Tao](#), undergraduate student at WHU, 2022, now PhD student at UW-Madison.
6. [Tian Bian](#), PhD student at CUHK, 2021, now PhD student at CUHK.
5. [Gabriel Gozum](#), undergraduate student at UW-Madison, 2021, now perception engineer at Applied Intuition.
4. [Eric Wang](#), undergraduate student at UW-Madison, 2021, now research assistant at UCSD.

3. Zhenxi Zhu, undergraduate student at BUPT, 2019, now PhD student at Nanjing University.
2. [Alex Singh](#), undergraduate student at CMU, 2019, now AI Engineer at Tesla.
1. [Pengna Li](#), undergraduate student at XJTU, 2019, now PhD student at XJTU.

Service

Reviewer for

- NeurIPS, ICML, ICLR
- CVPR, ECCV, ICCV
- WACV, AAAI, IJCAI, MICCAI
- TMLR, Nature Communications, IJCV, TCSVT, TMM, TIP

Additional Information

Language skills

- Native speakers of Mandarin with professional English speaking capability .

Programming skills

- Proficient with Python, TensorFlow, PyTorch. Familiar with C, Matlab, C++.